
Muhammad Yunas

[Speaker Bio](#) • [Slack Intro](#) • [Newsletter Bio](#) • [Predictive FAQ](#)

Brand Reference Document • Physical AI & Industrial IoT GTM • Prepared May 2026 • West Islip, NY

Brand Identity Pillars

Physical AI & Industrial IoT GTM Leadership • Agentic AI Commercialization • OT/IT Convergence

- Industry 4.0 Category Creation • Predictive Operations • Cyber-Physical Resilience • Industrial Autonomy

ARR Built	Team Scale	Quota Attainment	Enterprise Accounts	F500 Partners
\$40M+	18 People	137%	12	6
ARR engine delivered	GTM team led	Quota performance	Enterprise wins	Fortune 500 partners

SECTION 1

Speaker Bio

Three versions for different use contexts — full conference bio, short program bio, and social media profile. Copy and paste the version appropriate for your platform or event.

Version A — Full Speaker Bio

~400 WORDS • [CONFERENCE KEYNOTES](#) • [SPEAKING AGENCY PROFILES](#) • [EVENT WEBSITES](#)

Muhammad Yunas is a category-defining Physical AI and Industrial IoT GTM executive operating at the exact intersection of Agentic AI, OT/IT convergence, and industrial revenue architecture. He is one of a rare breed of commercial leaders who can walk a plant floor at 6 AM and close a Fortune 500 boardroom deal by afternoon — fluent in the language of both operational technology and enterprise procurement.

Muhammad has spent his career building the commercial infrastructure that turns breakthrough industrial technology into scalable revenue. He has constructed \$40M+ ARR engines, led 18-person cross-functional GTM teams, and launched Physical AI categories from zero to market — driving 137% quota attainment across 12 enterprise accounts and forging strategic partnerships with 6 Fortune 500 companies. His work spans manufacturing, energy, oil & gas, and smart factory environments where the cost of getting AI wrong is measured not in lost subscriptions, but in lost production, compromised safety, and operational failure.

His expertise sits at the convergence of five critical disciplines: Agentic AI commercialization, OT/IT convergence strategy, Industry 4.0 market development, predictive operations, and cyber-physical resilience. Muhammad is known for his ability to translate complex AI and IoT architectures — edge inference engines, digital twins, industrial data fabrics, autonomous control loops — into compelling commercial narratives that move enterprise buyers from pilot hesitancy to committed deployment.

Muhammad is an active voice in the Physical AI conversation. He is the author of *Silicon Meets Steel*, a landmark article on the commercialization of Physical AI in industrial environments. He contributes to *Board Memo* on the topic of Industrial Autonomy, and has published GTM case studies documenting the frameworks he has built and deployed. His writing reflects a practitioner's perspective: grounded in data, honest about friction, and relentlessly focused on what actually works at scale.

Muhammad speaks on the following topics:

- Physical AI GTM Strategy: Building the Commercial Layer for Industrial Intelligence
- OT/IT Convergence Revenue Architecture: Selling Across the Operational Divide
- Agentic AI in Industrial Operations: From Lab to Plant Floor at Enterprise Scale
- Cyber-Physical Resilience as a Board Imperative: Why Industrial Security Is a Revenue Conversation

- Industry 4.0 Category Creation: How to Build a Market That Doesn't Exist Yet

Muhammad Yunas is available for keynotes, panel discussions, and executive roundtables on Physical AI, Industrial IoT, and the future of industrial autonomy.

Version B — Short Speaker Bio

~150 WORDS • EVENT PROGRAMS • CONFERENCE AGENDAS • SESSION INTRODUCTIONS

Muhammad Yunas is a Physical AI and Industrial IoT GTM leader who builds the commercial layer of industrial intelligence companies. With \$40M+ ARR delivered, 18-person GTM teams led, and 137% quota attainment across 12 enterprise accounts and 6 Fortune 500 partnerships, Muhammad occupies a rare position: equal parts plant floor operator and enterprise revenue architect.

He specializes in Agentic AI commercialization, OT/IT convergence strategy, Industry 4.0 category creation, and cyber-physical resilience — and is recognized for translating the most complex industrial AI platforms into scalable, boardroom-ready commercial narratives. Muhammad is the author of *Silicon Meets Steel* on Physical AI and a contributor to *Board Memo* on Industrial Autonomy.

He is available for keynotes, panels, and executive roundtables on Physical AI, Industrial IoT, and the future of industrial autonomy.

Version C — Twitter / X Bio

280 CHARACTERS MAX • TWITTER / X PROFILE • SOCIAL MEDIA HEADER

Muhammad Yunas | Physical AI & Industrial IoT GTM Leader | Agentic AI | OT/IT Convergence |
Industry 4.0 | Building the commercial layer of industrial autonomy |
[linkedin.com/in/muhammadyunas](https://www.linkedin.com/in/muhammadyunas)

Character count: 196 / 280

SECTION 2

Slack / Community Intro

~250 WORDS • SLACK WORKSPACES • DISCORD CHANNELS • ONLINE COMMUNITIES • AI / IIOT / GTM FORUMS

Hey everyone — excited to join this community. A quick intro:

I'm **Muhammad Yunas**, a Physical AI and Industrial IoT GTM leader based in West Islip, NY. I build the commercial layer of Physical AI and Industrial IoT companies — GTM strategy, category creation, revenue architecture, and partner ecosystems. My background spans manufacturing, energy, and smart factory environments, where I've built \$40M+ ARR engines, led 18-person GTM teams, and driven 137% quota attainment across enterprise and Fortune 500 accounts.

My professional obsessions: **Agentic AI in industrial environments**, OT/IT convergence, predictive operations, and why cyber-physical resilience needs to be a board conversation — not just an IT ticket. I write about this in *Silicon Meets Steel* and contribute to *Board Memo* on Industrial Autonomy. The gap between AI that works in a demo and AI that works reliably on a plant floor is wide, expensive, and mostly a GTM problem — that's the problem I solve.

Right now I'm **exploring VP GTM / CRO opportunities** at Physical AI and industrial autonomy companies shaping the next decade of manufacturing, energy, and smart factory operations. If your company is building in this space — or you know someone who is — I'd love to connect.

Happy to talk about go-to-market strategy for industrial tech, Agentic AI commercialization, or how to align OT and IT buyers in a complex enterprise deal. DM me here or find me on LinkedIn:

[linkedin.com/in/muhammadyunas](https://www.linkedin.com/in/muhammadyunas)

Looking forward to learning from and contributing to this community.

SECTION 3

Newsletter / Publication Bio

~100 WORDS • NEWSLETTER BYLINES • PODCAST SHOW NOTES • PUBLICATION AUTHOR BOXES • PRESS RELEASES

Muhammad Yunas is a Physical AI and Industrial IoT GTM leader specializing in Agentic AI commercialization, OT/IT convergence strategy, and Industry 4.0 revenue architecture. With \$40M+ ARR built across manufacturing and energy verticals, he is recognized for translating complex industrial AI platforms into scalable commercial engines. Muhammad writes and speaks on Physical AI, industrial autonomy, and cyber-physical resilience. Connect at linkedin.com/in/muhammadyunas.

SECTION 4

Predictive FAQ / Executive Reference

Twenty fully answered questions organized by category — written in Muhammad's first-person voice for use in executive conversations, recruiter calls, investor meetings, and board-level discussions. Answers are substantive, data-backed, and ready for immediate use.

Category A: Career

&

Positioning

Q1. What makes you uniquely positioned for a Physical AI GTM leadership role?

Most GTM leaders in this space come from one of two directions: they understand enterprise software sales, or they understand industrial operations. I've lived both. I've built \$40M+ ARR engines in markets where the buyer is a VP of Engineering who hasn't trusted a software vendor in a decade — and I've done it by understanding what's actually happening on the plant floor, not just what's in the pitch deck. I speak OT. I understand why a SCADA engineer cares more about uptime than features, why a plant manager's buying decision is fundamentally different from an enterprise IT leader's, and why most AI vendors fail in industrial markets not because their technology is wrong, but because their commercial motion is wrong. I also bring the enterprise discipline — pipeline rigor, partner leverage, category positioning, quota accountability — that early-stage Physical AI companies desperately need as they try to scale beyond pilots. The combination is rare. I've built it over a career at the intersection of the two, and it shows in the results: 137% quota, 12 enterprise accounts, 6 F500 partnerships.

Q2. How do you define Physical AI, and why does it matter commercially?

Physical AI is artificial intelligence that perceives, reasons about, and acts on the physical world in real time — not just digital data. It's the convergence of edge computing, industrial sensors, computer vision, predictive analytics, autonomous control systems, and Agentic AI workflows embedded directly into physical operational environments: factories, energy grids, pipelines, refineries, ports. Commercially, it matters because the stakes are categorically different from enterprise SaaS. A Physical AI system that underperforms doesn't just reduce productivity — it can shut down a production line,

compromise worker safety, or create a cyber-physical security incident with regulatory and liability consequences. That raises the buying bar enormously, but it also raises the ROI ceiling. When Physical AI works, it's delivering tens of millions of dollars in recovered yield, avoided downtime, and optimized energy consumption annually. That's not a productivity tool — it's a strategic asset. The commercial opportunity is massive, the competitive window is open right now, and most vendors are still trying to sell it like SaaS. That's the gap I work in.

Q3. What's the difference between your GTM approach and a traditional enterprise SaaS GTM leader?

Three fundamental differences. First, my ICP definition goes deeper than firmographic and technographic data — I map the operational technology environment, the SCADA stack, the age of existing PLCs, the maintenance philosophy, and the union and regulatory constraints that shape how change gets adopted on a plant floor. SaaS GTM rarely goes that far. Second, my sales motion accounts for the OT/IT alignment problem. In industrial deals, IT controls budget and cybersecurity requirements, while OT controls operational requirements and adoption. They often don't agree, and a deal can die in the gap between them. I've built playbooks specifically for navigating that dynamic. Third, my success metrics include operational outcomes — uptime improvement, yield recovery, energy efficiency — not just ARR and NRR. Enterprise SaaS GTM optimizes for renewal and expansion within a software relationship. Physical AI GTM optimizes for operational proof, which then drives enterprise commitment. The commercial motion is slower and more complex, but the contracts are larger and the churn is near zero when you get

the deployment right.

Q4. What's the largest GTM program you've led, and what were the outcomes?

The largest program I've led was a full-category GTM buildout for a Physical AI and Industrial IoT platform entering manufacturing and energy verticals with no established market presence. I came in at the beginning — no ICP, no sales playbook, no partner infrastructure, no category language the market could recognize. I built the ICP framework, defined the category narrative around OT/IT convergence and predictive operations, recruited and developed an 18-person GTM team across sales, solution engineering, and partner development, and established a partner ecosystem with 6 Fortune 500 strategic partners. I also built the sales enablement infrastructure: competitive positioning, technical proof-of-concept frameworks, OT buyer playbooks, and executive briefing materials. The outcome was \$40M+ ARR, 12 enterprise accounts closed, and 137% of quota achieved. More importantly, we created a repeatable commercial motion — not just a set of individual deals — that the organization could scale. That's what I consider the real measure of GTM success at this level.

Q5. Why are you exploring new opportunities now?

I'm at a stage in my career where I want to apply everything I've built — the category creation playbooks, the OT/IT alignment methodology, the enterprise

revenue architecture — to a company where it can have maximum impact from day one. The Physical AI and industrial autonomy market is entering what I'd call its commercial inflection point. The technology has matured enough to deliver proof, the enterprise appetite is real, and the companies that build commercial infrastructure now will define the category for the next decade. I want to be at the helm of that buildout — at the VP GTM or CRO level — at a company that's serious about winning the industrial market, not just piloting in it. I bring a proven playbook, a clear-eyed understanding of what industrial buyers actually need, and the track record to de-risk the investment. The timing, from my perspective, is exactly right.

Category B: GTM Strategy

&

Methodology

Q6. How do you approach category creation in a new Physical AI or Industrial IoT market?

Category creation in Physical AI starts with a naming problem: if your buyers don't have a word for what you're selling, they can't budget for it, and they can't champion it internally. The first thing I do is establish the category language — specific, operational, and rooted in outcomes buyers already care about. "Predictive operations platform" lands differently than "AI for manufacturing." Once the language exists, I build the thought leadership infrastructure: articles, executive briefings, conference positioning, analyst

relationships. The goal is to make the company synonymous with the category before competitors realize the category exists. Simultaneously, I identify the lighthouse accounts — the 2 or 3 enterprises whose logos and operational outcomes will define what's possible in the market. Everything else gets resourced around those early wins. The mistake most companies make is trying to do broad market development and enterprise landing simultaneously. I sequence them: define the category, land the lighthouse, then scale the motion.

Q7. Walk me through your ICP definition process for an industrial technology company.

My ICP process for industrial technology is significantly more granular than a standard B2B ICP framework. I start with vertical selection — not all manufacturing is the same, and discrete manufacturing, process manufacturing, energy generation, and oil & gas each have fundamentally different operational technology environments, regulatory constraints, and buying behaviors. Within the selected vertical, I profile by asset intensity, operational complexity, and OT infrastructure maturity. A greenfield smart factory with a modern data historian is a completely different buyer than a 40-year-old refinery running legacy SCADA. From there, I map the buying committee: who owns the operational outcome (usually a VP of Operations or Plant Manager), who owns the technology decision (VP of IT or CTO), who owns the security requirement (CISO), and who controls capital allocation (CFO or COO). Each persona has different

success criteria. The ICP isn't complete until I've documented all four and built messaging that speaks to each one's specific risk and reward profile. Then I validate the ICP in the field — with lost deals as much as won ones — and iterate quarterly.

Q8. How do you align OT buyers and IT buyers in a complex enterprise deal?

This is the central challenge of every significant Physical AI deal, and most vendors fail at it because they try to run one sales motion for both audiences. I run two parallel tracks with explicit coordination. The OT track is outcome-first: I lead with operational proof — reduced downtime, improved yield, predictive maintenance ROI — and I make sure the OT stakeholder sees technology that integrates with their existing operational environment without requiring a rip-and-replace. I never make an OT buyer feel like IT is taking over their floor. The IT track is risk-first: I lead with security architecture, data governance, integration standards, and compliance alignment. IT needs to know that what we're deploying doesn't create a new attack surface or a shadow-IT problem. The coordination point is a joint executive briefing — typically at the 60-day mark — where both stakeholders see a unified operational and technical case. I also designate an internal champion from each side who can translate between the two domains. When both sides feel heard and protected, deals close. When one feels like they're losing the conversation, deals stall or die.

Q9. What's your philosophy on partner ecosystem development in industrial markets?

In industrial markets, the partner ecosystem isn't optional — it's load-bearing. Industrial enterprises don't buy from vendors they don't trust, and trust in OT environments is built over years of operational relationships. The right partners — systems integrators, OEM equipment manufacturers, industrial automation vendors, and strategic technology alliances — bring credibility, access, and operational context that no direct sales team can replicate. My philosophy is to treat partners as co-creators of value, not channels for distribution. That means investing in joint solution development, co-selling with shared pipeline accountability, and building partner-specific enablement that makes them genuinely competent to sell and implement the solution — not just refer leads. With my 6 Fortune 500 partner relationships, I built joint go-to-market programs with shared revenue targets, co-authored technical integration playbooks, and created customer success frameworks that defined how value delivery was shared. Partnerships in industrial markets succeed when both parties win on the outcome, not just on the transaction.

Q10. How do you build a sales enablement program for a team selling Physical AI to OT environments?

Standard sales enablement fails in OT environments because it focuses on features and competitive positioning — the things that matter in a software evaluation — without addressing the operational context that determines whether a Physical AI solution gets adopted on a plant floor. My enablement programs are built in three layers. The first layer is domain fluency: every seller on my team learns the OT stack — SCADA, PLCs, DCS, historians, edge gateways — well enough to have a credible conversation with a controls

engineer. Not deep expertise, but enough not to embarrass themselves or lose a deal on a basic vocabulary gap. The second layer is buyer psychology: understanding how OT buyers evaluate risk, how they relate to vendors, and what signals tell them a vendor is serious about their operational environment versus just selling software into it. The third layer is proof architecture: how to build and present a proof of concept that maps directly to the buyer's operational KPIs, so the path from pilot to production contract is clear and defensible. Enablement that covers all three layers produces sellers who win in OT environments consistently, not occasionally.

Category C: Industry Expertise

Q11. How do you define OT/IT convergence, and why is it a strategic imperative?

OT/IT convergence is the integration of operational technology — the systems that run physical industrial processes — with information technology — the systems that manage data, compute, and enterprise applications. Historically, these two domains were air-gapped by design: OT was isolated because connectivity meant vulnerability, and IT operated in a separate organizational and technical world. That separation is ending. The economic pressure to extract real-time operational intelligence from industrial assets — to enable predictive maintenance, autonomous optimization, energy efficiency, and supply chain responsiveness — requires connecting OT environments to enterprise data infrastructure and cloud compute. The convergence is a

strategic imperative for two reasons. First, it's the enabling condition for Physical AI: you cannot deploy Agentic AI in a factory if you can't get reliable, real-time data from the plant floor into a compute environment where AI can act on it. Second, it's a significant security risk if done without a deliberate architecture — connecting OT to IT exposes critical physical infrastructure to cyber threats that previously couldn't reach it. Companies that manage the convergence intelligently create durable competitive advantage. Companies that ignore it, or rush it, create existential operational risk.

Q12. What is Agentic AI in an industrial context, and what are the highest-ROI use cases?

Agentic AI in industrial contexts refers to AI systems that don't just analyze data and surface recommendations — they perceive operational conditions, reason about optimal responses, and take autonomous action within defined operational parameters, continuously and in real time. This is distinct from predictive analytics, which tells an operator what might happen; Agentic AI acts on what's happening without waiting for human intervention in the loop. The highest-ROI use cases I've seen are: autonomous predictive maintenance — where the system detects equipment degradation signatures, schedules maintenance interventions, and dispatches work orders without human initiation, recovering millions in avoided unplanned downtime; autonomous process optimization — where the AI continuously adjusts process parameters in response to raw material variability, environmental conditions, or demand shifts, improving yield and energy efficiency simultaneously; and autonomous quality inspection — where computer vision and edge inference replace manual inspection at production rates no human team can match, reducing

defect escape rates and warranty costs. The common thread is that these use cases replace not just human labor, but human reaction time — and in industrial environments, reaction time is often the difference between a manageable event and a costly failure.

Q13. What does cyber-physical resilience mean, and why is it a board-level concern?

Cyber-physical resilience is the capacity of an industrial organization to maintain safe, continuous physical operations in the face of cyber threats targeting the systems that control those operations. It's not just cybersecurity — it's the intersection of cybersecurity, operational continuity, and physical safety. When a ransomware attack hits an IT system, the consequence is typically data loss or business disruption. When a cyber attack hits an OT system — a SCADA network, a distributed control system, an industrial IoT platform — the consequence can be a pipeline shutdown, a power grid failure, a chemical release, or a manufacturing line that produces defective products at scale without detection. Those consequences are measured in human safety, regulatory liability, insurance exposure, and reputational damage — not just lost revenue. That's why it's a board-level concern: it sits at the intersection of operational risk, financial risk, and reputational risk simultaneously. As Physical AI and OT/IT convergence expand the attack surface of industrial environments, boards that don't have a cyber-physical resilience framework are making an unexamined bet on a risk that can be existential. I've built the business case for this conversation at the board level, and I know how to make it land.

Q14. What are the key differences in selling to manufacturing vs. energy vs. oil & gas verticals?

The differences are real and consequential — treating these as a single "industrial" vertical is a mistake that kills deals. In manufacturing, the primary value driver is yield, throughput, and quality — buyers are measuring ROI in production output per hour. The buying cycle is often faster because the operational proof is visible within a production quarter. The OT environment is more variable, ranging from legacy discrete manufacturing to highly automated continuous process environments. In energy — utilities, renewables, grid management — the primary drivers are uptime, regulatory compliance, and increasingly, carbon efficiency and grid resilience. Buyers are extremely risk-averse because their regulatory exposure is high and their operational complexity is enormous. Procurement cycles are long and involve multiple layers of regulatory review. In oil

&

gas, the dominant concerns are safety, environmental compliance, and asset integrity management — the consequence of failure is measured in human lives and environmental liability, not just lost production. Buyers in this vertical are deeply skeptical of new technology until it has been proven in comparable operating conditions. The sales motion requires patience, deep domain credibility, and a willingness to invest in proof before the commercial conversation matures. I've sold successfully in all three, and I adapt the commercial motion to the vertical's specific risk and reward architecture.

Q15. How do you see the Physical AI market evolving over the next 3–5 years?

We are at the very beginning of a 10-to-15-year industrial transformation that Physical AI will drive. In the next 3 to 5 years, I expect three major dynamics. First, consolidation of the platform layer: the fragmented landscape of point solutions — predictive maintenance tools, computer vision systems, industrial analytics platforms — will consolidate into integrated Physical AI platforms that provide a unified operational intelligence and autonomous action layer. The companies that win this consolidation will be those that build the most trusted OT data infrastructure and the most capable Agentic AI layer on top of it. Second, the OT security imperative will become a buying prerequisite: cyber-physical resilience will move from a differentiator to a table-stakes requirement as regulatory frameworks in critical infrastructure tighten globally. This will advantage companies with purpose-built OT security architecture. Third, the talent gap will accelerate GTM differentiation: there are very few people who can sell Physical AI credibly at an enterprise level, and the companies that build that commercial capability fastest will win disproportionate market share. That's precisely why I believe the commercial buildout — the GTM layer — is the most strategically important investment an industrial AI company can make right now.

Category D: Leadership

&

Team Building

Q16. How do you build a world-class Physical AI GTM team from scratch?

I've done this, and the principles I've learned are counterintuitive. The first hire is not a seller — it's a solution engineer with deep OT domain expertise. Nothing kills Physical AI GTM faster than a sales team that can't survive a technical conversation with a controls engineer. Once I have technical credibility anchored, I hire commercially aggressive enterprise sellers with industrial sector experience — people who are comfortable in a plant environment, not just a conference room. I build in parallel: direct sellers who can run complex enterprise deals independently, and partner development leaders who can build and manage the ecosystem relationships that give us market reach beyond our direct team's capacity. I also invest early in sales operations — pipeline hygiene, forecasting rigor, and win/loss analysis — because in an emerging category, understanding why you're losing is as commercially valuable as winning. Over the 18-person team I built, I maintained a deliberate ratio of OT domain expertise to enterprise commercial experience, and I built a culture where both domains respected and learned from each other. That mutual respect is rare and genuinely differentiating.

Q17. How do you manage the tension between OT domain expertise and commercial sales talent?

This tension is real and it's one of the defining management challenges in Physical AI GTM. OT domain experts often distrust salespeople — they've seen too many vendors oversell and underdeliver in industrial environments. Commercial sellers often underestimate the complexity of OT environments and try to run their standard enterprise motion into a context it wasn't designed for. I manage this by making the shared outcome explicit: we all win when a

customer deploys successfully and gets the operational results we promised. That's the forcing function that aligns both camps. Structurally, I pair OT solution engineers and commercial sellers into account teams with shared quota — when they're measured on the same outcome, they find ways to work together. I also invest in cross-training: sellers learn enough OT to be credible, and domain experts learn enough commercial to understand why pipeline velocity and deal qualification matter. I've found that the best Physical AI commercial professionals are people who started with genuine curiosity about the operational environment and developed commercial discipline on top of that foundation — not the other way around.

Q18. What's your approach to revenue forecasting in an emerging technology category?

Forecasting in an emerging category is genuinely harder than forecasting in a mature market, and the worst thing you can do is pretend otherwise. Standard CRM-based probability weighting fails when you're selling something buyers have never bought before, because deal velocity is unpredictable and buyer hesitation patterns are different from established technology categories. My approach starts with stage-gating rigor: I define very specific, observable buyer behaviors that advance a deal from stage to stage, so forecast accuracy depends on what the buyer has actually done — signed an NDA, committed an internal champion, approved a proof-of-concept budget — not on how optimistic the seller feels about the relationship. I also maintain a separate track for category maturation indicators — inbound inquiry volume, conference engagement, analyst coverage — because in an emerging category, market-level signals are as important as individual deal signals for accurate revenue

prediction. And I'm disciplined about separating demand generation pipeline from partner-sourced pipeline, because they have different velocity profiles. Finally, I build the board forecast with a base case, upside case, and category breakthrough scenario — because in a new market, outlier events — a major competitor validating the category, a high-profile customer win going public — can change revenue trajectory in ways a single-number forecast can't capture.

Q19. How do you maintain quota attainment discipline during a category creation phase?

Category creation is not a license to miss quota, and I've never allowed it to become one on my teams. The discipline I've applied starts with honest quota setting: I set quota targets that reflect the actual market development stage, not the aspirational ARR the company wishes it had. Aggressive but achievable is the frame — and I push back hard on board-level quota expectations that ignore market maturity realities, because a demoralized team chasing impossible numbers in an emerging market does more damage than a realistic quota with clear accountability. Within the quota framework, I maintain deal-level accountability through weekly pipeline reviews that focus on buyer action, not seller activity. I distinguish between deals that are progressing because buyers are engaged and deals that are progressing because sellers are optimistic. I also create early-harvest targets within the category — smaller proof-of-concept contracts, departmental pilots, and partner-sourced deals — that generate recognizable revenue and quota credit while the larger enterprise contracts develop through their longer cycle. The 137% quota attainment I've delivered reflects the combination of realistic target-setting and disciplined pipeline management — not an easy market or a

favorable macro environment.

Q20. What's one thing most companies get wrong about GTM for Physical AI?

They build a software GTM motion and aim it at an industrial market. I see this consistently: a company with genuinely powerful Physical AI technology hires a team of enterprise SaaS veterans, gives them a standard enterprise GTM playbook, and sends them into manufacturing, energy, or oil

&

gas accounts. Twelve months later, they have a pipeline full of stalled pilots and no production deployments — and they conclude either that the technology isn't ready or the market isn't ready. Neither is usually true. What isn't ready is the commercial motion. Industrial buyers don't buy the way enterprise software buyers buy. They don't evaluate on a feature checklist, they don't sign 30-day trials, and they don't trust vendors who can't speak their operational language. The entire GTM architecture — the ICP, the discovery motion, the proof-of-concept structure, the commercial terms, the success metrics, the expansion strategy — needs to be rebuilt for the industrial buyer's decision environment. That rebuild is exactly what I do. And when it's done right, the deals are larger, the retention is higher, and the customer relationships are deeper than anything I've seen in enterprise SaaS. The market is extraordinary — it just requires a commercial layer that actually fits it.

Document Reference & Contact

Muhammad Yunas • Physical AI & Industrial IoT GTM Leader • West Islip, NY

LinkedIn: [linkedin.com/in/muhammadyunas](https://www.linkedin.com/in/muhammadyunas)

Authored Works: *Silicon Meets Steel* (Physical AI) • Board Memo (Industrial Autonomy) • GTM Case Studies

Available For: Keynotes • Panel Discussions • Executive Roundtables • VP GTM / CRO Opportunities in Physical AI & Industrial Autonomy

Prepared May 2026 • Muhammad Yunas Brand Reference Document • All content written for immediate use. Copy-paste ready.